


Honors Data Structures

Lecture 21: Shortest Paths.

4/11 & 4/13
2022

Daniel Bauer



Google

Maps

Start address e.g., "SFO"

End address e.g., "94526"

seattle

sydney

Get Directions

[Search the map](#)
[Find businesses](#)
[Get directions](#)

[Search Results](#)
[My Maps](#)

[Print](#)
[Send](#)
[Link to this page](#)

3.

Merge onto **I-5 N** via the ramp on the **left** to **Vancouver BC**

1.0 mi
4.

Take exit **167** on the **left** toward **Seattle Center**

0.7 mi
5.

Turn **right** at **Fairview Ave N**

400 ft
6.

Turn **left** at **Valley St**

0.2 mi
7.

Turn **right** at **Westlake Ave N**

1.6 mi
8.

Turn **right** at **4th Ave N**

0.3 mi
9.

Turn **right** at **N 34th St**

0.3 mi
10.

Turn **right** at **Stone Way N**

115 ft
11.

Turn **left** at **N Northlake Way**

0.3 mi
12.

Kayak across the **Pacific Ocean**
Entering Australia (New South Wales)

7,906 mi
13.

Sharp **right** at **Macquarie St**

0.4 mi
14.

Turn **right** at **Albert St**

292 ft
15.

Turn **left** at **Phillip St**

0.1 mi
16.

Turn **right** at **Bridge St**

0.3 mi
17.

Turn **left** at **George St**

0.2 mi

B

To: Sydney NSW Australia

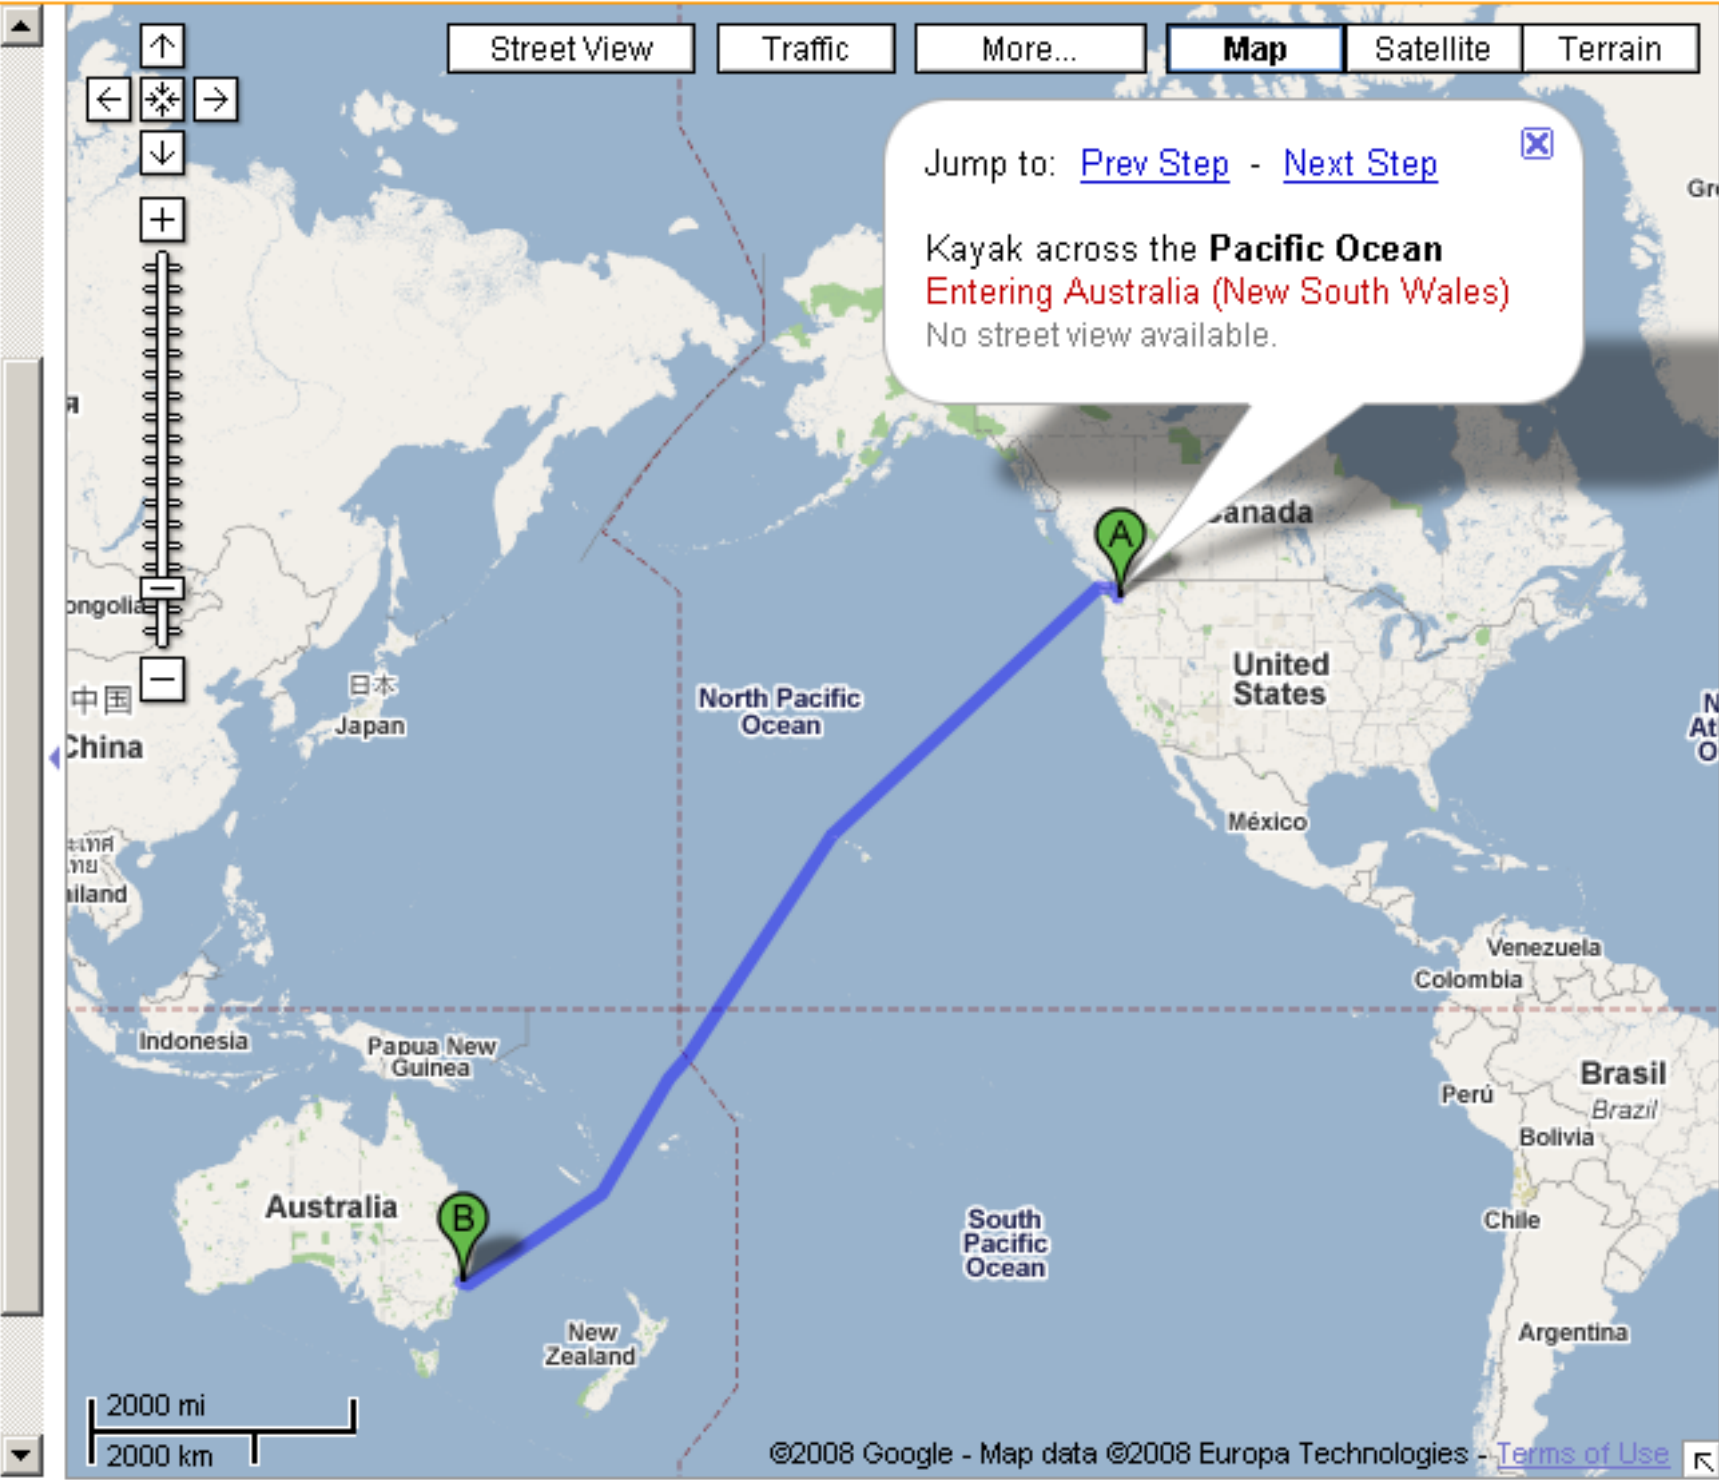
✕

Edit

[km](#) | [miles](#)

[Add destination ...](#)

These directions are for planning purposes only. You may find that




Graph Traversals

- Different ways of exploring graphs:
 - Topological sort for Directed Acyclic Graphs.
 - Depth First Search (a generalization of pre-order traversal on trees to graphs, uses a Stack)
 - Breadth First Search (uses a Queue)
 - Dijkstra's algorithm to find weighted shortest paths (uses a Priority Queue)

Finding Shortest Paths


- Goal: Find the shortest path between two vertices s and t .



What is the shortest path between v_3 and v_7 ?

Finding Shortest Paths


- Goal: Find the shortest path between two vertices s and t .



What is the shortest path between v_3 and v_7 ?


Finding Shortest Paths

- Goal: Find the shortest path between two vertices s and t .
- It turns out that finding the shortest path between s and ALL other vertices is just as easy. This problem is called **single-source shortest paths**.




Finding Shortest Paths

- Goal: Find the shortest path between two vertices s and t .
- It turns out that finding the shortest path between s and ALL other vertices is just as easy. This problem is called **single-source shortest paths**.




Finding Shortest Paths

- Goal: Find the shortest path between two vertices s and t .
- It turns out that finding the shortest path between s and ALL other vertices is just as easy. This problem is called **single-source shortest paths**.




Finding Shortest Paths

- Goal: Find the shortest path between two vertices s and t .
- It turns out that finding the shortest path between s and ALL other vertices is just as easy. This problem is called **single-source shortest paths**.



Breadth First Search



Queue: { V_1 }

Queue **q**

q.enqueue(start)


while (**q** is not empty):

u = q.dequeue()

 for each **v** adjacent to **u**:

q.enqueue(v)

Breadth First Search



Queue: { V₁ }

Queue **q**

q.enqueue(start)

Set **discovered**

while (**q** is not empty):

u = q.dequeue()


 for each **v** adjacent to **u**:

 if (**v** is not in **discovered**):

q.enqueue(v)

discovered.add(v)

Breadth First Search



Queue: {}

Queue **q**

q.enqueue(start)

Set **visited**

while (**q** is not empty):

u = q.dequeue()


 for each **v** adjacent to **u**:

 if (**v** is not in **visited**):

q.enqueue(v)

visited.add(v)

Breadth First Search



Queue: { V₂, V₄ }

Queue **q**

q.enqueue(start)

Set **visited**

while (**q** is not empty):

u = q.dequeue()


 for each **v** adjacent to **u**:

 if (**v** is not in **visited**):

q.enqueue(v)

visited.add(v)

Breadth First Search



Queue: $\{V_4, V_5\}$

Queue **q**

q.enqueue(start)

Set **visited**

while (**q** is not empty):

u = q.dequeue()


 for each **v** adjacent to **u**:

 if (**v** is not in **visited**):

q.enqueue(v)

visited.add(v)

Breadth First Search



Queue: $\{V_5, V_3, V_6, V_7\}$

Queue **q**

q.enqueue(start)

Set **visited**

while (**q** is not empty):

u = q.dequeue()


 for each **v** adjacent to **u**:

 if (**v** is not in **visited**):

q.enqueue(v)

visited.add(v)

Breadth First Search



Queue: $\{V_3, V_6, V_7\}$

Queue **q**

q.enqueue(start)

Set **visited**

while (**q** is not empty):

u = q.dequeue()


 for each **v** adjacent to **u**:

 if (**v** is not in **visited**):

q.enqueue(v)

visited.add(v)

Breadth First Search



Queue: $\{V_6, V_7\}$

Queue **q**

q.enqueue(start)

Set **visited**

while (**q** is not empty):

u = q.dequeue()


 for each **v** adjacent to **u**:

 if (**v** is not in **visited**):

q.enqueue(v)

visited.add(v)

Breadth First Search



Queue: {V₇}

Queue **q**

q.enqueue(start)

Set **visited**

while (**q** is not empty):

u = q.dequeue()


 for each **v** adjacent to **u**:

 if (**v** is not in **visited**):

q.enqueue(v)

visited.add(v)

Breadth First Search



Queue: {}

Queue **q**

q.enqueue(start)

Set **visited**

while (**q** is not empty):

u = q.dequeue()


 for each **v** adjacent to **u**:

 if (**v** is not in **visited**):

q.enqueue(v)

visited.add(v)

Breadth First Search



Queue **q**

q.enqueue(start)

Set **visited**

while (**q** is not empty):

u = q.dequeue()

 for each **v** adjacent to **u**:


 if (**v** is not in **visited**):

q.enqueue(v)

visited.add(v)

Running time: $O(|V|+|E|)$

Breadth First Search



Running time: $O(|V|+|E|)$

Queue **q**

q.enqueue(start)


Set **visited**

while (**q** is not empty):

BFS will traverse the entire graph even without a visited set.

What happens if we use a stack (DFS)?

Unweighted Shortest Paths



Queue: $\{V_1\}$


for all v :

$v.cost = \infty$

$v.prev = \text{null}$

start.cost = 0

Unweighted Shortest Paths



Queue: { V₁ }

for all **v**:

v.cost = ∞


v.prev = **null**

start.cost = 0

Queue **q**

q.enqueue(start)

Unweighted Shortest Paths



Queue: {}

for all **v**:

v.cost = ∞

v.prev = **null**

start.cost = 0

Queue **q**

q.enqueue(start)

while (**q** is not empty):

u = **q.dequeue()**

 for each **v** adjacent to **u**:


 if **v.cost** == ∞ :

v.cost = **u.cost** + 1

v.prev = **u**

q.enqueue(v)

Unweighted Shortest Paths



Queue: { V₂, V₄ }

for all **v**:

v.cost = ∞

v.prev = **null**

start.cost = 0

Queue **q**

q.enqueue(start)

while (**q** is not empty):

u = **q.dequeue()**

 for each **v** adjacent to **u**:

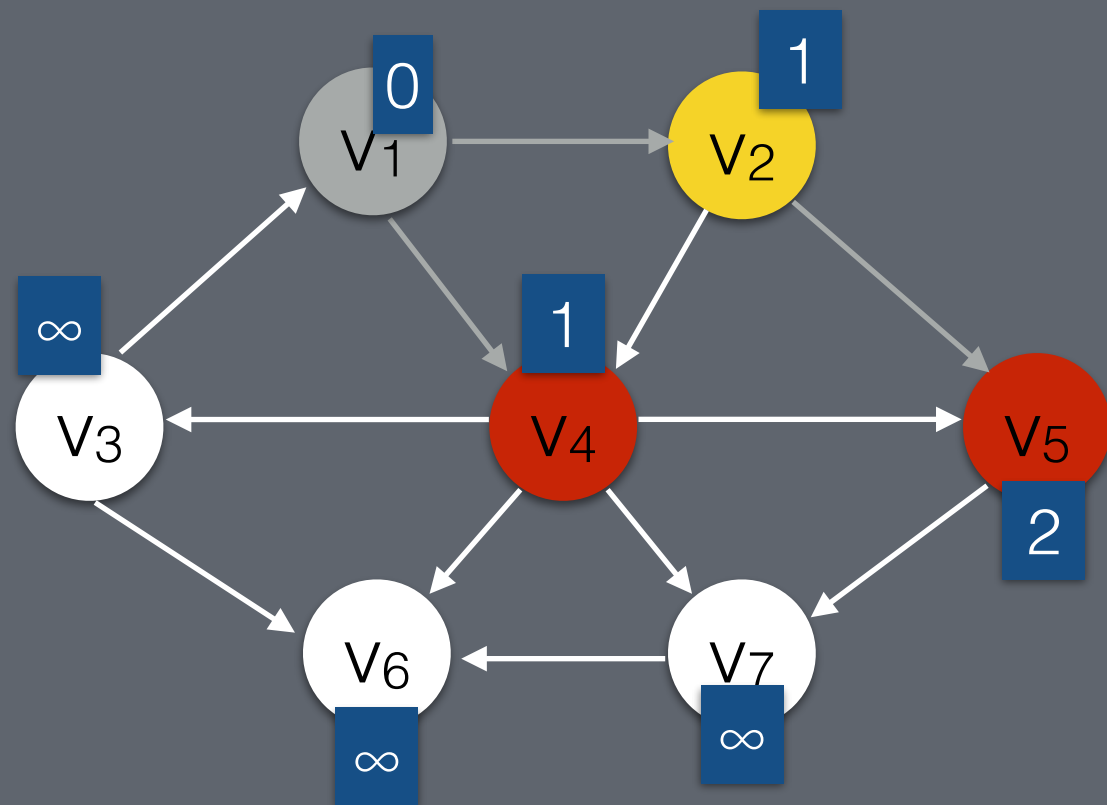
 if **v.cost** == ∞:

v.cost = **u.cost** + 1

v.prev = **u**

q.enqueue(v)

Unweighted Shortest Paths



Queue: { V₄, V₅ }

for all **v**:

v.cost = ∞

v.prev = **null**

start.cost = 0

Queue **q**

q.enqueue(start)

while (**q** is not empty):

u = **q.dequeue()**

 for each **v** adjacent to **u**:

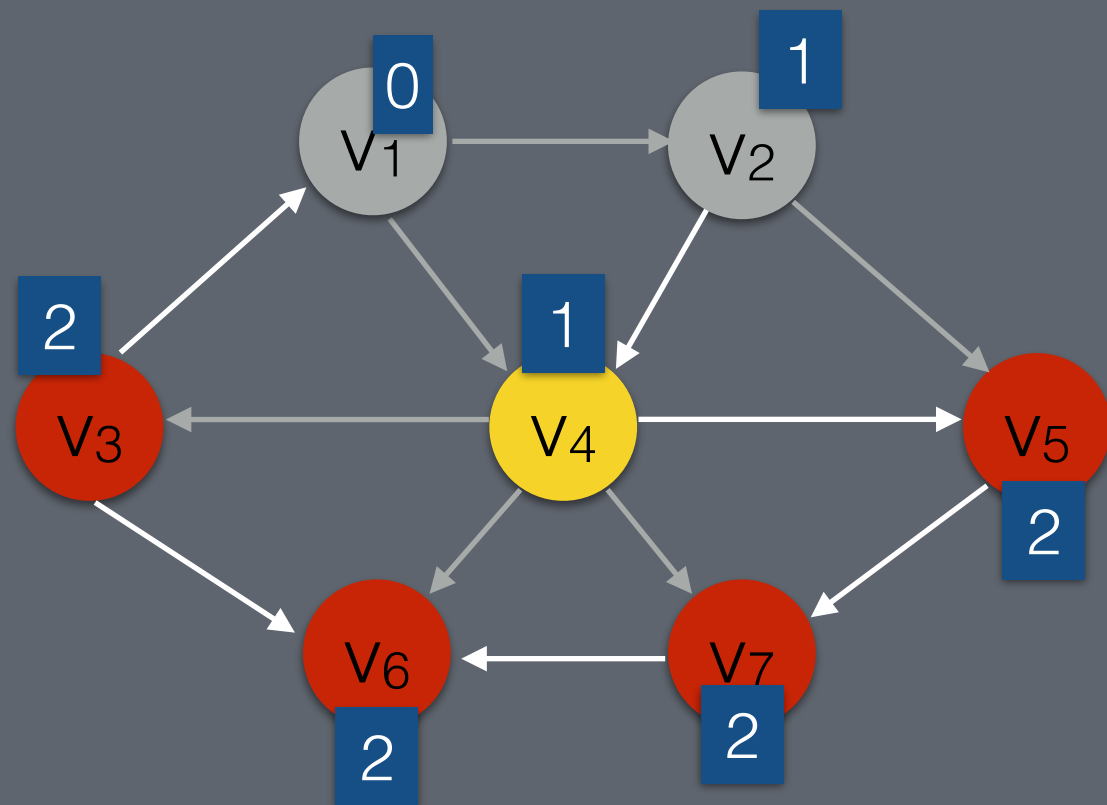
 if **v.cost** == ∞ :

v.cost = **u.cost** + 1

v.prev = **u**

q.enqueue(v)

Unweighted Shortest Paths



Queue: $\{V_5, V_3, V_6, V_7\}$

for all **v**:

v.cost = ∞

v.prev = **null**

start.cost = 0

Queue **q**

q.enqueue(**start**)

while (**q** is not empty):

u = **q**.dequeue()

 for each **v** adjacent to **u**:

 if **v**.cost == ∞ :


v.cost = **u**.cost + 1

v.prev = **u**

q.enqueue(**v**)

Weighted Shortest Paths


- Goal: Find the shortest path between two vertices s and t .



What is the shortest path between v_2 and v_6 ?

Weighted Shortest Paths

- Goal: Find the shortest path between two vertices s and t .
- Normal BFS will find this path.




length 2
cost 11

What is the shortest path between v_2 and v_6 ?

Weighted Shortest Paths

- Goal: Find the shortest path between two vertices s and t .
- This path has a lower cost.




length 3
cost 8

What is the shortest path between v_2 and v_6 ?

Negative Weights

- We normally expect the shortest path to be simple.
- Edges with Negative Weights can lead to negative cycles.
- The concept of “shortest path” is then not clearly defined.



What is the shortest path between v_2 and v_6 ?

Dijkstra's Algorithm for Weighted Shortest Path

Dijkstra's Algorithm for Weighted Shortest Path

- Cost annotations for each vertex reflect the lowest cost *using only vertices visited so far*.

Dijkstra's Algorithm for Weighted Shortest Path

- Cost annotations for each vertex reflect the lowest cost *using only vertices visited so far*.
- There might be a lower-cost path through other vertices that have not been seen yet.


Dijkstra's Algorithm for Weighted Shortest Path

- Cost annotations for each vertex reflect the lowest cost *using only vertices visited so far*.
 - There might be a lower-cost path through other vertices that have not been seen yet.
- Keep nodes on a **priority queue** and always expand the vertex with the lowest cost annotation first! ← This is a **greedy** algorithm

Dijkstra's Algorithm for Weighted Shortest Path

- Cost annotations for each vertex reflect the lowest cost *using only vertices visited so far*.
 - There might be a lower-cost path through other vertices that have not been seen yet.
- Keep nodes on a **priority queue** and always expand the vertex with the lowest cost annotation first! ← This is a **greedy** algorithm
- Intuitively, this means we will never overestimate the cost and miss lower-cost path.

Dijkstra's Algorithm



for all **v**:

v.cost = ∞

v.visited = **false**


v.prev = **null**

start.cost = 0

PriorityQueue **q**

q.insert(start)

Dijkstra's Algorithm



for all **v**:

v.cost = ∞

v.visited = **false**

v.prev = **null**

start.cost = 0

PriorityQueue **q**

q.insert(start)

while (**q** is not empty):

u = **q.pollMin()**

u.visited = **true**

visit vertex **u**

for each **v** adjacent to **u**:

if **not v.visited**:

c = **u.cost** + **cost(u,v)**

if (**c** < **v.cost**):


v.cost = **c**

v.prev = **u**

q.insert(v)

discover and
relax vertex **v**

Dijkstra's Algorithm



for all **v**:

v.cost = ∞

v.visited = **false**

v.prev = **null**

start.cost = 0

PriorityQueue **q**

q.insert(start)

while (**q** is not empty):

u = **q.pollMin()**

u.visited = **true**

visit vertex **u**

for each **v** adjacent to **u**:

if **not v.visited**:

c = **u.cost** + **cost(u,v)**

if (**c** < **v.cost**):


v.cost = **c**

v.prev = **u**

q.insert(v)

discover and
relax vertex **v**

Dijkstra's Algorithm



for all **v**:

v.cost = ∞

v.visited = **false**

v.prev = **null**

start.cost = 0

PriorityQueue **q**

q.insert(start)

while (**q** is not empty):

u = **q.pollMin()**

u.visited = **true**

visit vertex **u**

for each **v** adjacent to **u**:

if **not v.visited**:

c = **u.cost** + **cost(u,v)**

if (**c** < **v.cost**):


v.cost = **c**

v.prev = **u**

q.insert(v)

discover and
relax vertex **v**

Dijkstra's Algorithm



for all **v**:

v.cost = ∞

v.visited = **false**

v.prev = **null**

start.cost = 0

PriorityQueue **q**

q.insert(start)

while (**q** is not empty):

u = **q.pollMin()**

u.visited = **true**

visit vertex **u**

for each **v** adjacent to **u**:

if **not v.visited**:

c = **u.cost** + **cost(u,v)**

if (**c** < **v.cost**):


v.cost = **c**

v.prev = **u**

q.insert(v)

discover and
relax vertex **v**

Dijkstra's Algorithm



for all **v**:

v.cost = ∞

v.visited = **false**

v.prev = **null**

start.cost = 0

PriorityQueue **q**

q.insert(start)

while (**q** is not empty):

u = **q.pollMin()**

u.visited = **true**

visit vertex **u**

for each **v** adjacent to **u**:

if **not v.visited**:

c = **u.cost** + **cost(u,v)**

if (**c** < **v.cost**):


v.cost = **c**

v.prev = **u**

q.insert(v)

discover and
relax vertex **v**

Dijkstra's Algorithm



for all **v**:

v.cost = ∞

v.visited = **false**

v.prev = **null**

start.cost = 0

PriorityQueue **q**

q.insert(start)

while (**q** is not empty):

u = **q.pollMin()**

u.visited = **true**

visit vertex **u**

for each **v** adjacent to **u**:

if **not v.visited**:

c = **u.cost** + **cost(u,v)**

if (**c** < **v.cost**):

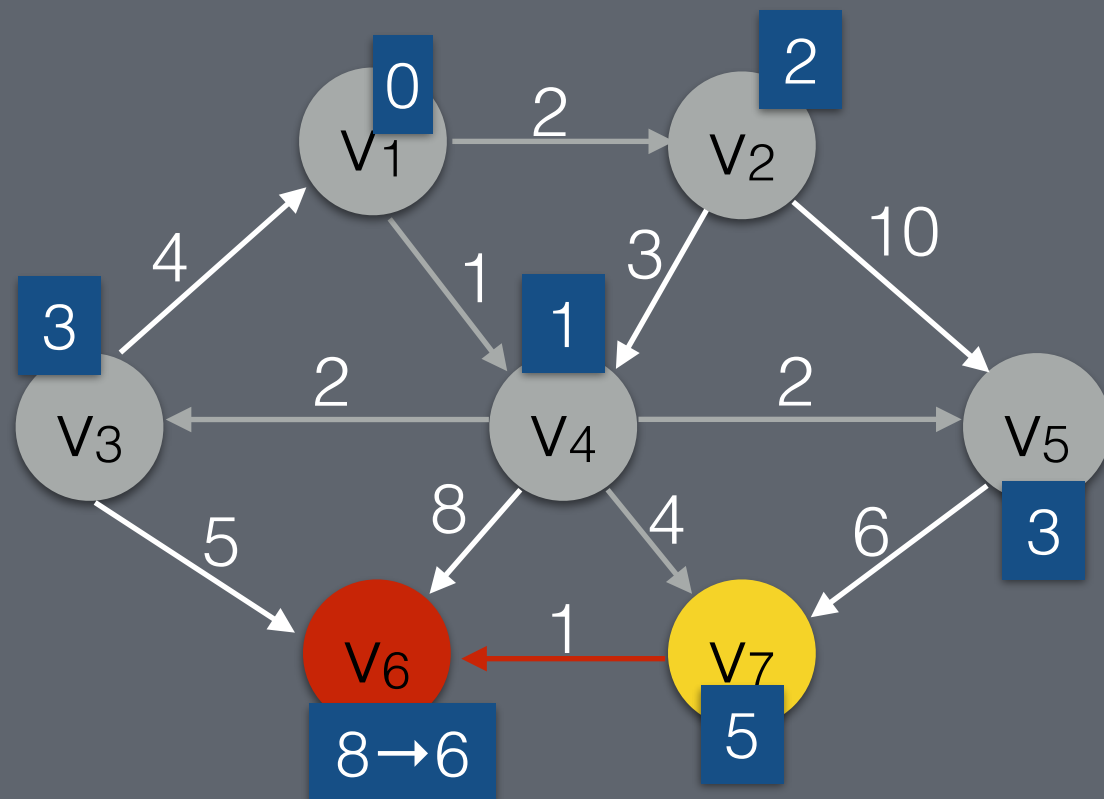
v.cost = **c**

v.prev = **u**

q.insert(v)

discover and
relax vertex **v**

Dijkstra's Algorithm



for all **v**:

v.cost = ∞

v.visited = **false**

v.prev = **null**

start.cost = 0

PriorityQueue **q**

q.insert(start)

while (**q** is not empty):

u = **q.pollMin()**

u.visited = **true**

visit vertex **u**

for each **v** adjacent to **u**:

if **not v.visited**:

c = **u.cost** + **cost(u,v)**

if (**c** < **v.cost**):


v.cost = **c**

v.prev = **u**

q.insert(v)

discover and
relax vertex **v**

Dijkstra's Algorithm



for all **v**:

v.cost = ∞

v.visited = **false**

v.prev = **null**

start.cost = 0

PriorityQueue **q**

q.insert(start)

while (**q** is not empty):

u = **q.pollMin()**

u.visited = **true**

visit vertex **u**

for each **v** adjacent to **u**:

if **not v.visited**:

c = **u.cost** + **cost(u,v)**

if (**c** < **v.cost**):

v.cost = **c**


v.prev = **u**

q.insert(v)

discover and
relax vertex **v**

Dijkstra's Algorithm

"Life Cycle" of a Vertex



Dijkstra's Running Time

- There are $|E|$ insert and deleteMin operations.
- The maximum size of the priority queue is $O(|E|)$. Each insert takes $O(\log |E|)$


$$\begin{aligned} &O(|E| \log |E|) \\ &= O(|E| \log |V|) \end{aligned}$$

$$\begin{aligned} &|E| \leq |V|^2, \text{ so} \\ &\log |E| \leq 2 \log |V| = O(\log |V|) \end{aligned}$$

A General View of Graph Search

Goals:


- Explore the graph systematically starting at s to
 - Find a vertex t / Find a path from s to t .
 - Find the shortest path from s to all vertices.
 - ...



A General View of Graph Search


In every step of the search we maintain

- The part of the graph already explored.
- The part of the graph not yet explored.
- A data structure (an agenda) of *next* edges (adjacent to the explored graph).



Agenda: (v_2, v_5) , (v_4, v_5) , (v_4, v_7)


A General View of Graph Search



Agenda: (v2,v5), (v4,v5), (v4,v7)

A General View of Graph Search

The graph search algorithms discussed so far differ almost only in the type of agenda they use:




Agenda: (v2,v5), (v4,v5), (v4,v7)

A General View of Graph Search

The graph search algorithms discussed so far differ almost only in the type of agenda they use:

- unweighted shortest paths: breadth first, uses a queue.




Agenda: (v2,v5), (v4,v5), (v4,v7)

A General View of Graph Search

The graph search algorithms discussed so far differ almost only in the type of agenda they use:

- unweighted shortest paths: breadth first, uses a queue.
- Dijkstra's: best first, uses a priority queue.




Agenda: (v2,v5), (v4,v5), (v4,v7)

A General View of Graph Search

The graph search algorithms discussed so far differ almost only in the type of agenda they use:


- unweighted shortest paths: breadth first, uses a queue.
- Dijkstra's: best first, uses a priority queue.
- Topological Sort: breadth first with constraint on items in the queue.



Agenda: (v2,v5), (v4,v5), (v4,v7)

Correctness of Dijkstra's Algorithm

- We want to show that Dijkstra's algorithm really finds the minimum path costs (we don't miss any shorter solutions by choosing the shortest edge greedily).
- Proof by induction on the set S of visited nodes.
- Base case:
 $|S|=1$. Trivial. Length shortest path is 0.



Correctness of Dijkstra's Inductive Step

- Assume the algorithm produces the minimal path cost from s for the subset S , $|S| = k$.
- Dijkstra's algorithm selects the next edge (u,v) leaving S .
- Assume there was a shorter path from s to v that does not contain (u,v) .
 - Then that path must contain another edge (x,y) leaving S .
 - The cost of (x,y) is already higher than (u,v) because we didn't choose it before (u,v)
- Therefore (u,v) must be on the shortest path.

